University of BRISTOL

1.2 Regular Expressions

Edwin Simpson

Department of Computer Science,

University of Bristol, UK.

bristol.ac.uk

How Can I Find Occurrences Of…

- Different word forms referring to the same thing:
  - *Woodchuck*, *woodchuck*, *woodchucks*, or *Woodchucks?*
  - *Barack Obama, Obama,* or *44th President of the US?*
  - *Happy, joyful, cheerful,* or *delighted?*

How Can I Find Occurrences Of…

- Phrases within a particular expression:
  - *The movie was very …* (e.g., *enjoyable/boring*)
  - *The patient's symptoms are …*

Regular Expressions (REs)

- A formal language for describing text strings
- Used in:
  - Unix tools like grep, emacs
  - Search
  - Computer languages
  - Parsing
  - Various NLP tasks and pipeline steps

Regular Expressions (REs)

- Used to **search for** phrases that match a pattern
  - Information extraction without machine learning!
  - Define text features for a classifier
- …and to **replace** occurrences of a pattern
  - Error correction
  - Preprocessing: use a more convenient form for downstream methods
  - E.g., so that *Woodchuck*, *woodchuck*, *woodchucks*, and *Woodchucks* are all replaced by a single form so they can be treated the same by a classifier

Regular Expressions (RegExps)

- RegExp: defines a set of strings
- RegExp search function: returns all occurrences of strings within that set from a corpus of text
- https://www.whatsmyip.org/regular-expression-tester/
- To search for “woodchuck”: /woodchuck/

Building Blocks for REs: Disjunctions

- Use **disjunctions** match with alternative characters
- E.g., */[Ww]oodchuck/*

Pattern	What does it match?
[Ww]	W or w
[A-Z]	Any upper-case letter
[a-z]	Any lower-case leter
[0-9]	Any single digit

Building Blocks: Wildcards, Negation, Repetition

Pattern	Role	Example	What does it match?
^ inside []	Negation	*[^A-Z]*	any char not an upper-case letter
.	Any character	*password.*	*password1, passwordx,…*
?	Optional	*woodchucks?*	*woodchuck, woodchucks*
*	Repeat 0 or more times	*Woodchucks**	*Woodchuck, Woodchucks, Woodchuckssssss, …*
+	Repeat 1 or more times	*Woo+dchucks*	*Woodchucks, Wooodchucks, Woooodchucks…*

Building Blocks: Combining Patterns

Pattern	Role	Example	What does it match?
|	Disjunction (logical OR)	*groundhog|woodchuck*	*groundhog, woodchuck*
()	Sub-expression	*wood(chuck|louse)*	*woodlouse, woodchuck*

Building Blocks: Start and End

Pattern	Role	Example	What does it match?
\b	Word boundary	*\bround*	*round* but not *ground*
^ outside []	Start of sentence	*^[Rr]*	*R* in *Round* or *r* in *round*
$	End of sentence	*.$*	*!* in *ground!*, *.* in *round.*

Constructing an RE: an Example

The word “the”

- Test it out at https://www.whatsmyip.org/regular-expression-tester/
- You can also try to write and RE for:
  - All strings with two consecutive repeated words, e.g., “the the”
  - All integers

Constructing an RE: an Example

The word “the”

- /the/

Constructing an RE: an Example

The word “the”

- /the/
- /[tT]he/ -- match capitalised/uncapitalised

Constructing an RE: an Example

The word “the”

- /the/
- /[tT]he/
- /\b[tT]he\b/ -- match only whole words

Constructing an RE: an Example

The word “the”

- /the/
- /[tT]he/
- /\b[tT]he\b/
- /[ˆa-zA-Z][tT]he[ˆa-zA-Z]/ -- match when ”the” is preceded or followed by non-letter tokens like ‘_’ or space

Constructing an RE: an Example

The word “the”

- /the/
- /[tT]he/
- /\b[tT]he\b/
- /[^a-zA-Z][tT]he[^a-zA-Z]/
- /(^|[^a-zA-Z])[tT]he([^a-zA-Z]|$)/ -- Also match when the word occurs at the start or end of a line

False Positives and Negatives

- Each line tries to reduce false +ves and –ves
  - False positive (FP) = a match that was incorrect
  - False negative (FN) = incorrectly did not match

	True label = +ve	True label = -ve
Prediction = +ve	**True positive (TP)**	**False positive (FP)**
Prediction = -ve	**False negative (FN)**	**True negative (TN)**

- NLP systems make a trade-off between these apparently opposing metrics

Summary

- REs encode text patterns that we want to find and replace
- Wide ranging uses in feature extraction and preprocessing
- Often used to implement a baseline model for many text mining tasks
- Referring to the introduction: REs are a tool for hard-coding knowledge
- For more complicated tasks, we will need to use machine learning

Reading

- Dan Jurafsky and James H. Martin. Speech and language processing (3rd edition draft). **Chapter 2.**
- https://web.stanford.edu/~jurafsky/slp3/